

# Lecture 3: Query processing and optimization

## Objectives:

- Reduce processing time
- Reduce buffer memory
- Reduce communication cost between sites
- Low resource usage

# Introduction

- *Functions of query processing:*
  - Transform a complicated query into a much simpler query.
  - This transformation must ensure **correctness** and **efficiency**
  - Each transformation method leads to different resource usages → **lowest resource usage**.

# Introduction

## *Simple transformation methods*

### Relational algebra:

- Simplify queries based on equivalent relational algebra expressions such that the querying time is minimized.
- This method disregards database structure and size.

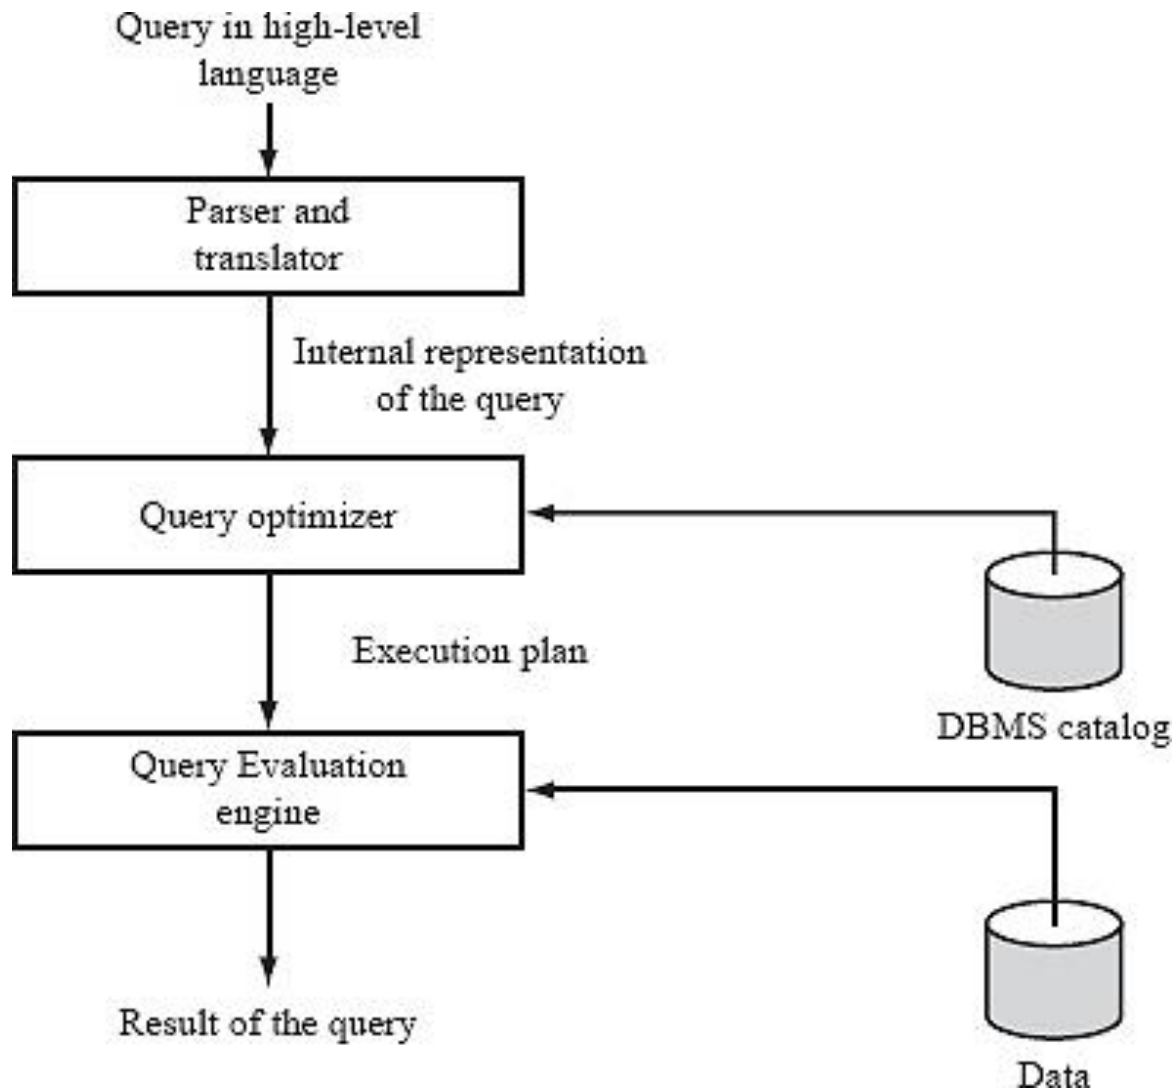
### Cost estimation:

- Specify data size and processing time of each element in the query.
- This method takes into account the data size and real processing time of the query.

# Querying Process

- Procedure of query execution
- Query pre-processing
- Query transformation
- Query optimization

# Querying Process



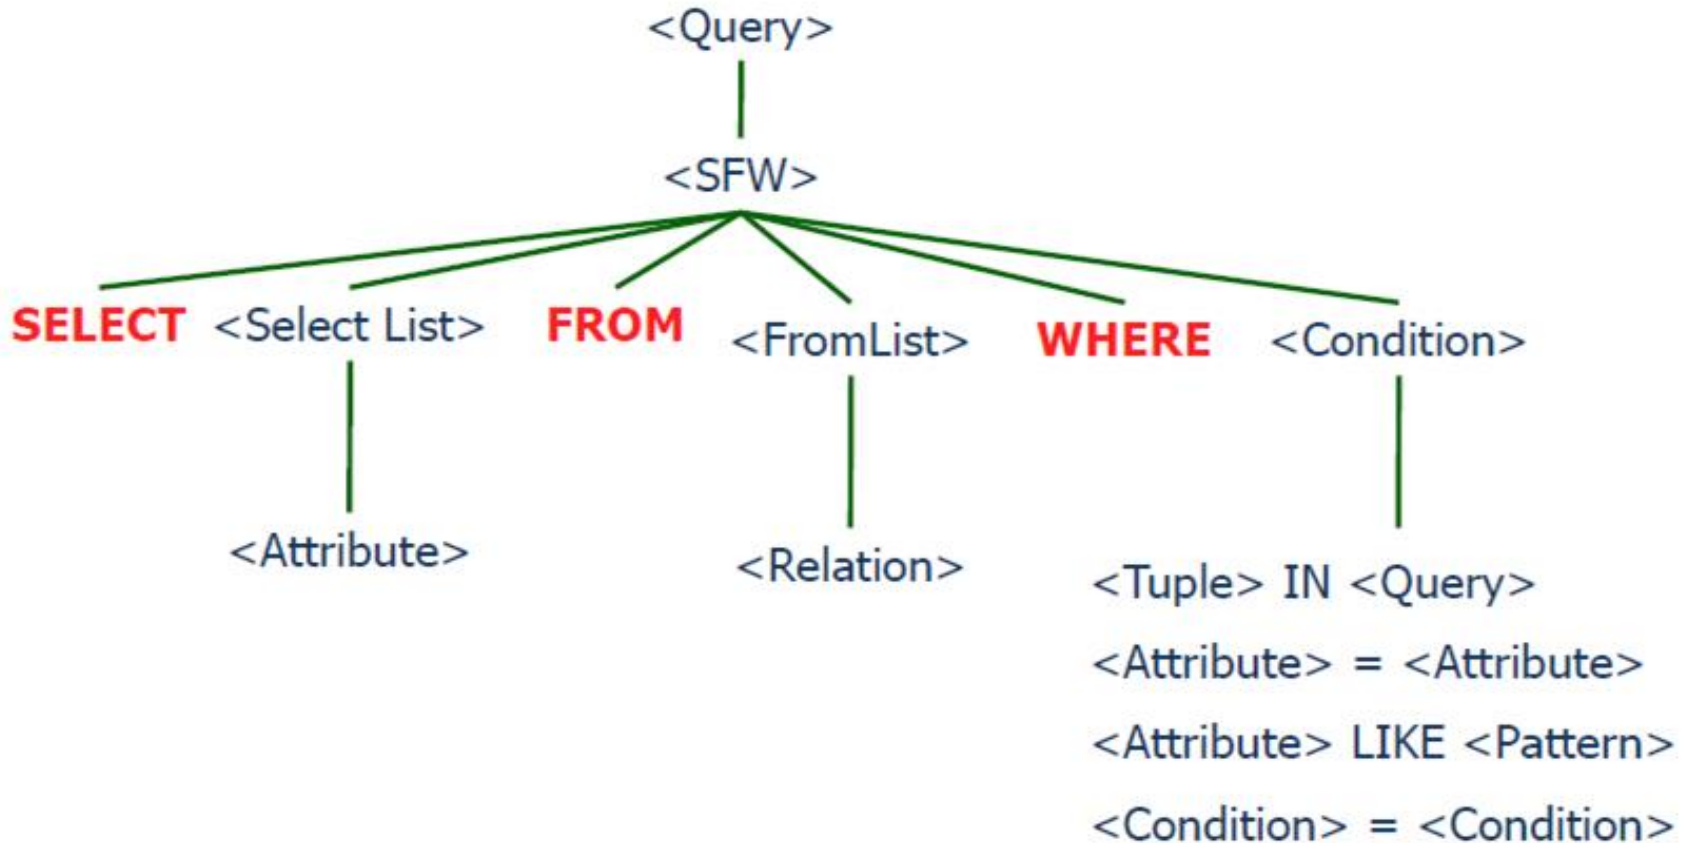
# Querying Process

- **Query preprocessor:**
  - Scanning: key words, attribute names, relations, ...
  - Parsing: syntax checking, parse tree representation
  - Validating: semantic checking (relations, attributes, data types)

# Querying Process

- **Query optimizer:**
  - Select suitable strategy for query processing
- **Query code generator:**
  - Generate codes to implement the plan
- **Runtime database processor:**
  - Compile codes to provide query results

# Querying Process





# Querying Process

Transform in to relational algebra expressions

- Query blocks: SELECT-FROM-WHERE-GROUP BY-HAVING
- Integrated queries: separate into query blocks

# Querying Process

Transform in to relational algebra expressions

SELECT Name, Salary

FROM Staff

WHERE Salary > (SELECT AVG(Salary)

FROM Staff

WHERE Gender = “Male”)

# Query Processing in Centralized DBS

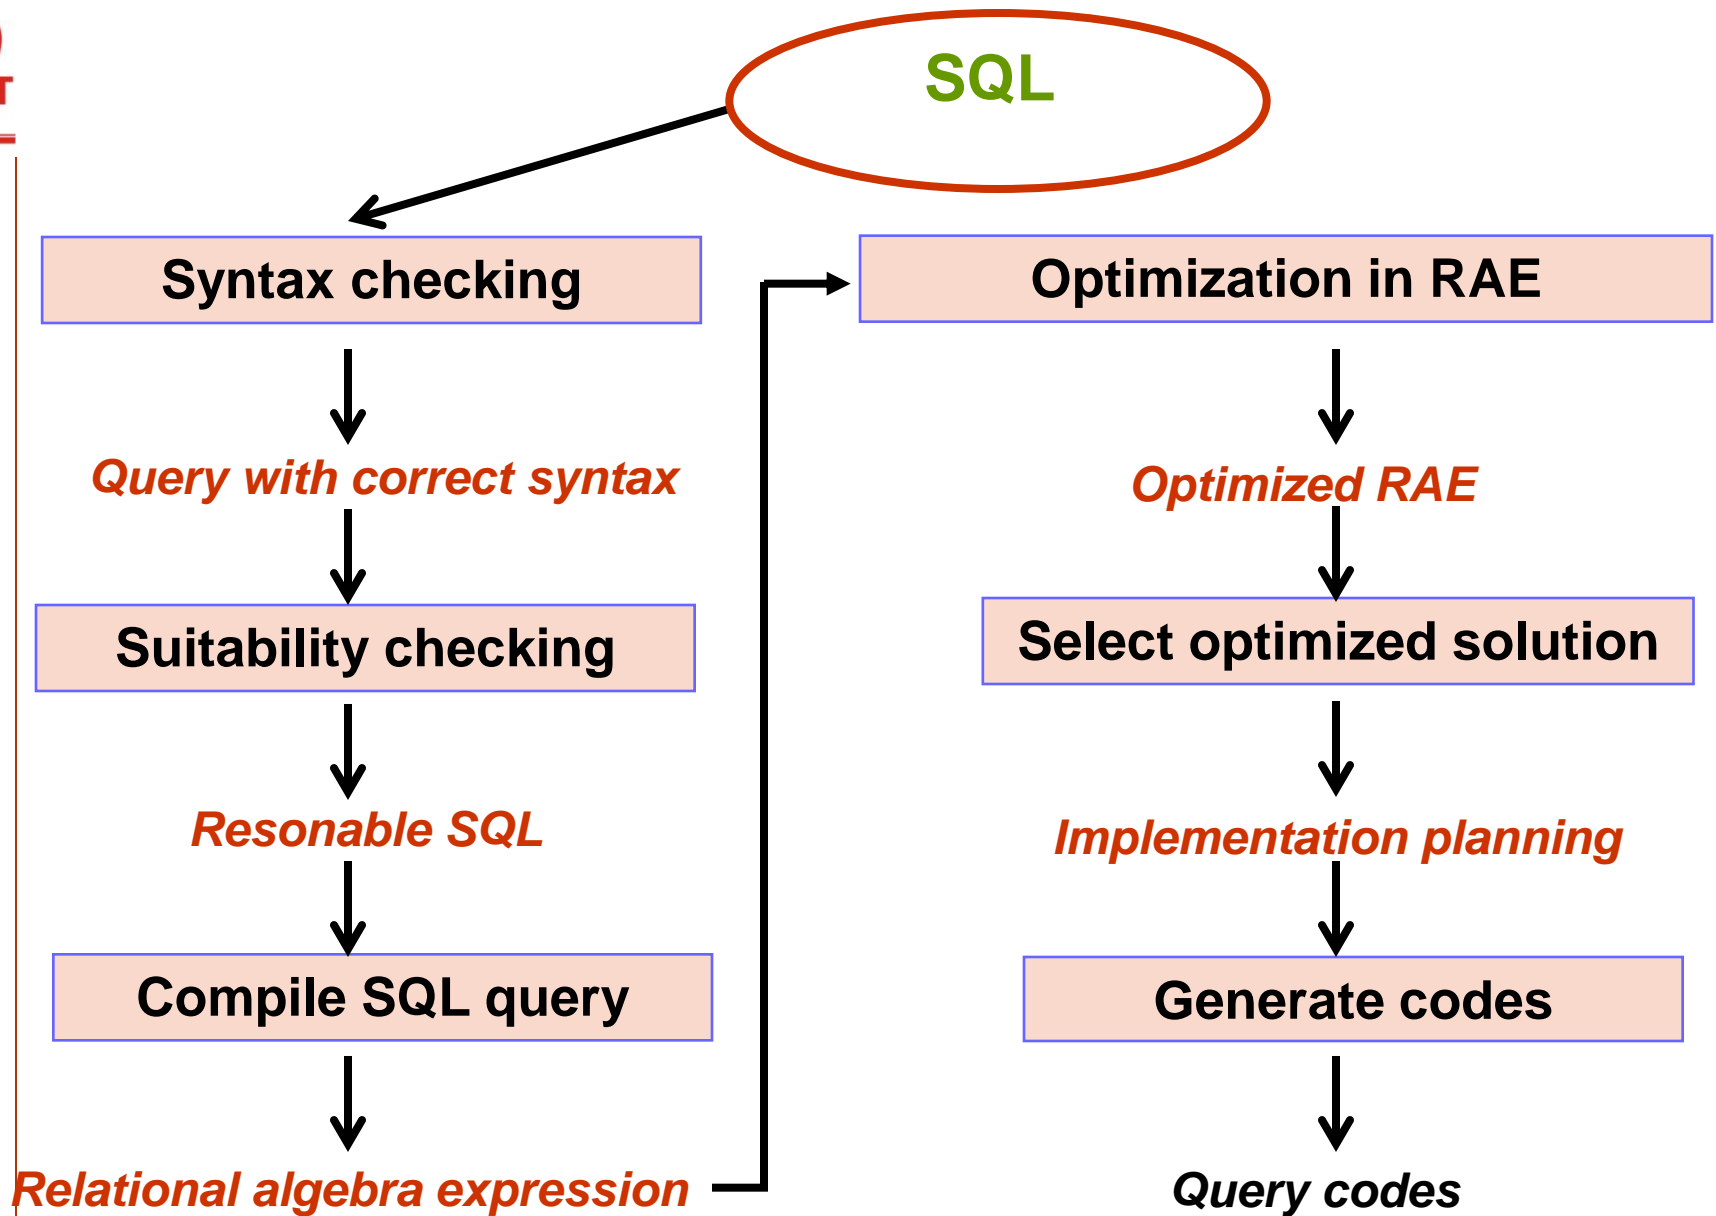
Comparison between centralized and distributed query processing:

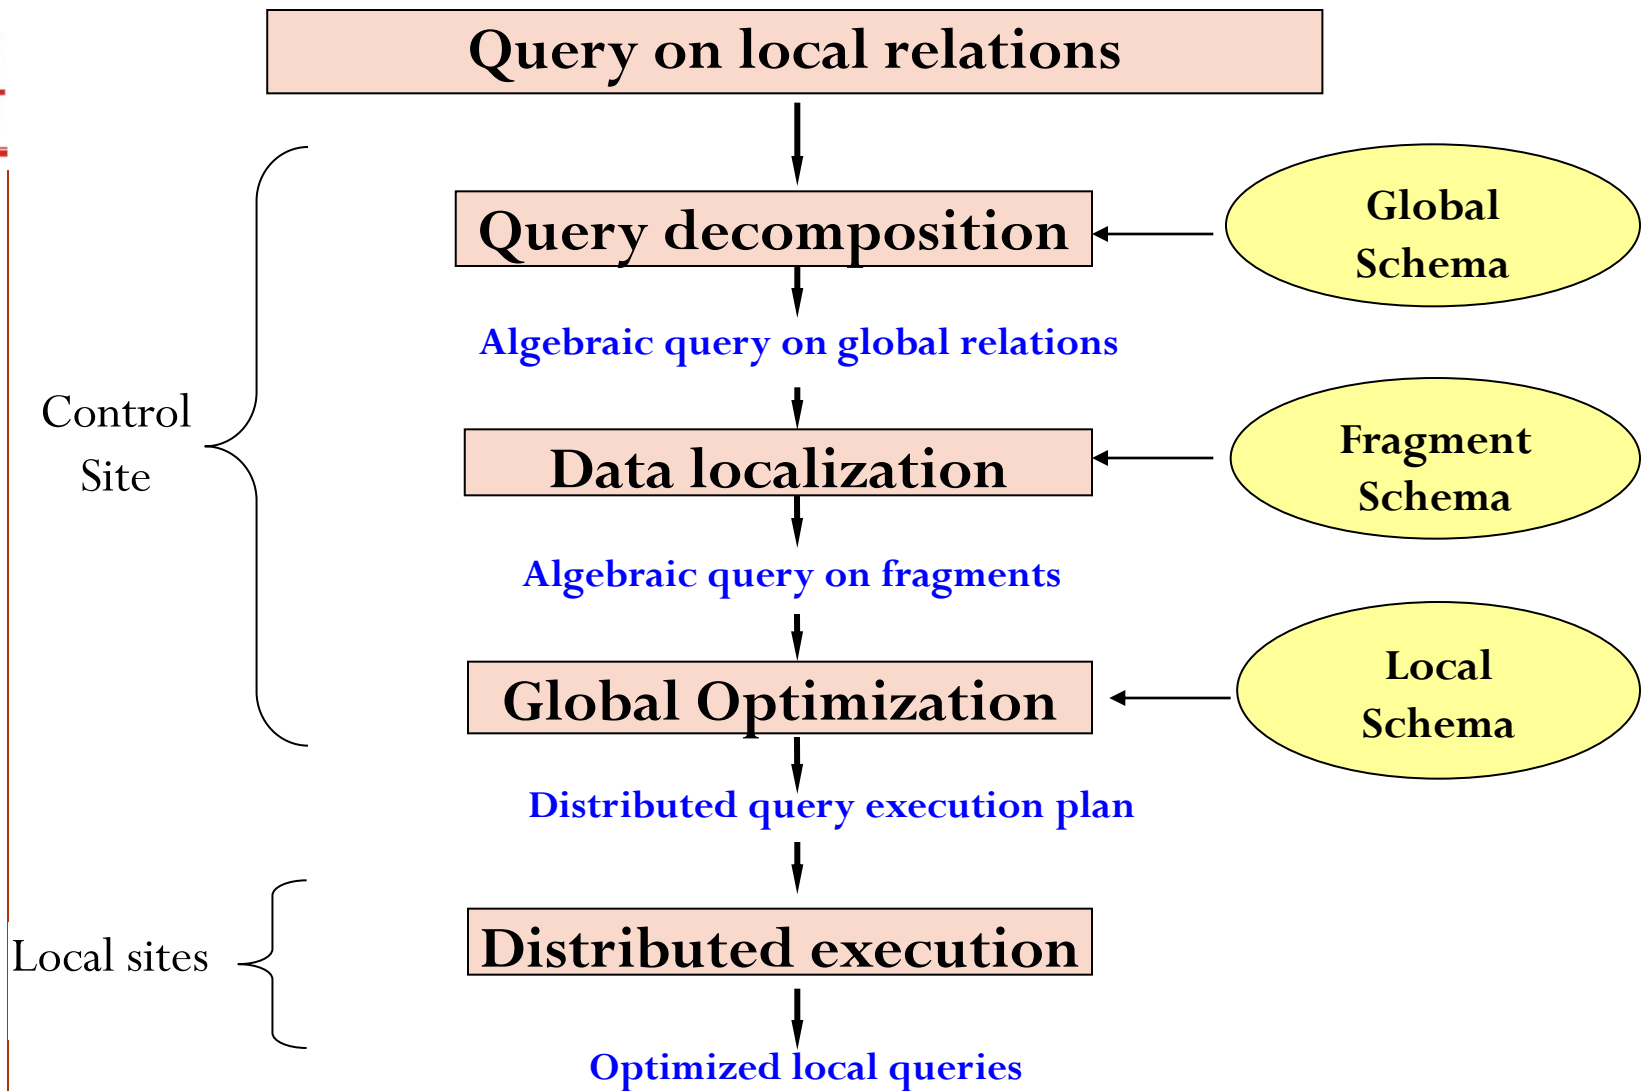
Centralized:

- Select the best relational algebra expression equivalent to the query
- Query processing strategies are described using relational algebra extensions

Distributed:

- Inherited from centralized environment
- More issues to concern:
  - Math expressions of data transmission between sites
  - Choose the best site to process data
  - Data transmission methods





Generic layering scheme for distributed query processing

# Query Processing in Centralized DBS

Optimization strategies in centralized databases:

- Distributed queries must be composed and processed in a centralized manner
- All distributed query optimization techniques are extensions from centralized query processing approaches
- Centralized query optimization is often simple

# Query Processing in Centralized DBS

INGRES algorithm: recursively breaks up a query expressed in SQL into smaller pieces which are executed along the way

- The query is first decomposed into a sequence of queries having a unique relation in common
- Each mono-relation query is processed by selecting, based on the predicate, the best access method to that relation

E.g: For example, if the predicate is of the form  $A = value$ , an index available on attribute A would be used if it exists.

However, if the predicate is of the form  $A \neq value$ , an index on A would not help, and sequential scan should be used.

# Query Processing in Centralized DBS

INGRES algorithm:

- Executes first the unary (monorelation) operations and tries to minimize the sizes of intermediate results in ordering binary (multirelation) operations
- Denote by  $q_{i-1} \rightarrow q_i$  a query  $q$  decomposed into two subqueries,  $q_{i-1}$  and  $q_i$ , where  $q_{i-1}$  is executed first and its result is consumed by  $q_i$
- Decomposes  $q$  into  $n$  subqueries  $q_1 \rightarrow q_2 \dots \rightarrow q_n$  using *detachment* and *substitution*.



# Query Processing in Centralized DBS

Detachment: breaks a query  $q$  into  $q' \rightarrow q''$ , based on a common relation that is the result of  $q'$

$q$ : **SELECT**  $R_2.A_2, R_3.A_3, \dots, R_n.A_n$   
**FROM**  $R_1, R_2, \dots, R_n$   
**WHERE**  $P_1(R_1.A_1)$  **AND**  $P_2(R_1.A_1, R_2.A_2, \dots, R_n.A_n)$ ;

$q'$ : **SELECT**  $R_1.A_1$  **INTO**  $R'_1$   
**FROM**  $R_1$   
**WHERE**  $P_1(R_1.A_1)$ ;

$q''$ : **SELECT**  $R_2.A_2, \dots, R_n.A_n$   
**FROM**  $R'_1, R_2, \dots, R_n$   
**WHERE**  $P_2(R_1.A_1, R_2.A_2, \dots, R_n.A_n)$ ;

## EMP (E)

ENO	ENAME	TITLE
A1	Nam	Phân tích HT
A2	Trung	Lập trình viên
A3	Đông	Phân tích HT
A4	Bắc	Phân tích HT
A5	Tây	Lập trình viên
A6	Hùng	Kỹ sư điện
A7	Dũng	Phân tích HT
A8	Chiến	Thiết kế DL

## ASG (G)

ENO	PNO	RESPONSIBILITY	DUR
A1	D1	Quản lý	12
A2	D1	Phân tích	34
A2	D2	Phân tích	6
A3	D3	Kỹ thuật	12
A3	D4	Lập trình	10
A4	D2	Quản lý	6
A5	D2	Quản lý	20
A6	D4	Kỹ thuật	36
A7	D3	Quản lý	48
A8	D3	Lập trình	15

## PRJ (J)

PNO	PNAME	BUDGET
D1	CSDL	20000
D2	CÀI ĐẶT	12000
D3	BẢO TRÌ	28000
D4	PHÁT TRIỂN	25000

## PAY(S)

TITLE	SALARY
Kỹ sư điện	1000
Phân tích HT	2500
Lập trình viên	3000
Thiết kế DL	4000

# Example

$q_1 =$  *“Find names of staffs working on project CSDL”*

$q_1$ :  
**SELECT** E.ENAME  
**FROM** E, G, J  
**WHERE** E.ENO = G.ENO **AND** G.PNO = J.PNO  
**AND** PNAME = “CSDL”;

$q_1$  is detached into  $q_{11} \rightarrow q'$ , where TGIAN1 is an intermediate relation.

$q_{11}$ :  
**SELECT** J.PNO **INTO** TEMP1  
**FROM** J  
**WHERE** PNAME = “CSDL”;

$q'$ :  
**SELECT** E.ENAME  
**FROM** E, G, TEMP1  
**WHERE** E.ENO = G.ENO  
**AND** G.PNO = TEMP1.PNO;

# Example

The successive detachments of  $q'$  may generate

$q_{12}$ : **SELECT** G.ENO **INTO** TEMP2  
**FROM** G,TEMP1  
**WHERE** G.PNO =TEMP1.PNO;

$q_{13}$ : **SELECT** E.ENAME  
**FROM** E,TEMP2  
**WHERE** E.ENO =TEMP2.ENO;

$q_1$  has been decomposed into  $q_{11} \rightarrow q_{12} \rightarrow q_{13}$

Query  $q_{11}$  is mono-relation and can be executed. However,  $q_{12}$  and  $q_{13}$  are not mono-relation and cannot be reduced by detachment.

# Query Processing in Centralized DBS

Tuple substitution: Given an  $n$ -relation query  $q$ , the tuples of one relation are substituted by their values, thereby producing a set of  $(n-1)$ -relation queries.

- One relation in  $q$  is chosen for tuple substitution. Let  $R_1$  be that relation.
- For each tuple  $t_{1i} \in R_1$ , the attributes referred to by in  $q$  are replaced by their actual values in  $t_{1i}$ , thereby generating a query  $q'$  with  $n-1$  relations.

$q(R_1, R_2, \dots, R_n)$  is replaced by  $\{q'(t_{1i}, R_2, R_3, \dots, R_n), t_{1i} \in R_1\}$

# Query Processing in Centralized DBS

Example: Let's consider the query  $q_{13}$

$q_{13}$ :       **SELECT**           E.ENAME  
               **FROM**           E, TEMP2  
               **WHERE**        E.ENO = TEMP2.ENO ;

The relation TEMP2 is over a single attribute (ENO). Assume that it contains only two tuples:  $\langle E1 \rangle$  and  $\langle E2 \rangle$ . The substitution of TEMP2 generates two one-relation subqueries:

$q_{131}$ :       **SELECT**           E.ENAME  
               **FROM**           E  
               **WHERE**        E.ENO = "E1";

$q_{132}$ :       **SELECT**           E.ENAME  
               **FROM**           E  
               **WHERE**        E.ENO = "E2";

These queries may then be executed

**Input:** MRQ: multirelation query with  $n$  relations

**Output:** Result of execution

**Begin**

Output  $\leftarrow \phi$

**If**  $n=1$  **then**

Output  $\leftarrow \text{run}(\text{MRQ})$

**Else**

$\{\text{ORQ}_1, \dots, \text{ORQ}_m, \text{MRQ}'\} \leftarrow \text{MRQ}$

**For**  $i \leftarrow 1$  **to**  $m$

Output'  $\leftarrow \text{run}(\text{ORQ}_i)$

Output  $\leftarrow \text{output} \cup \text{output}'$

**Endfor**

$R \leftarrow \text{CHOOSE\_RELATION}(\text{MRQ}')$

**For** each tuple  $t \in R$

$\text{MRQ}'' \leftarrow \text{substitute values for } t \text{ in } \text{MRQ}'$

Output'  $\leftarrow \text{INGRES-QOA}(\text{MRQ}'')$

Output  $\leftarrow \text{output} \cup \text{output}'$

**Endfor**

**Endif**

# Query Decomposition

Normalization: to transform the query to a normalized form to facilitate further processing.

- Conjunctive normal form:

$$(p_{11} \vee p_{12} \vee \dots \vee p_{1n}) \wedge \dots \wedge (p_{m1} \vee p_{m2} \vee \dots \vee p_{mn})$$

- Disjunctive normal form

$$(p_{11} \wedge p_{12} \wedge \dots \wedge p_{1n}) \vee \dots \vee (p_{m1} \wedge p_{m2} \wedge \dots \wedge p_{mn})$$

where  $p_{ij}$  is a simple predicate



# Query Decomposition

## Equivalence rules

- $p_1 \wedge p_2 \iff p_2 \wedge p_1$
- $p_1 \vee p_2 \iff p_2 \vee p_1$
- $p_1 \wedge (p_2 \wedge p_3) \iff (p_1 \wedge p_2) \wedge p_3$
- $p_1 \vee (p_2 \vee p_3) \iff (p_1 \vee p_2) \vee p_3$
- $p_1 \wedge (p_2 \vee p_3) \iff (p_1 \wedge p_2) \vee (p_1 \wedge p_3)$
- $p_1 \vee (p_2 \wedge p_3) \iff (p_1 \vee p_2) \wedge (p_1 \vee p_3)$
- $\neg(p_1 \vee p_2) \iff \neg p_2 \vee \neg p_1$
- $\neg(p_1 \wedge p_2) \iff \neg p_2 \wedge \neg p_1$
- $\neg(\neg p) \iff p$

# Query Decomposition

Example: “Find the names of employees who have been working on project P1 for 12 or 24 months”

```
SELECT      ENAME
FROM        EMP, ASG
WHERE       EMP.ENO = ASG.ENO
             AND    ASG.PNO = "P1"
             AND    (DUR = 12 OR DUR = 24);
```

conjunctive normal form is

$$\text{EMP.ENO} = \text{ASG.ENO} \wedge \text{ASG.PNO} = \text{"P1"} \wedge (\text{DUR} = 12 \vee \text{DUR} = 24)$$

disjunctive normal form is

$$(\text{EMP.ENO} = \text{ASG.ENO} \wedge \text{ASG.PNO} = \text{"P1"} \wedge \text{DUR} = 12) \vee$$
$$(\text{EMP.ENO} = \text{ASG.ENO} \wedge \text{ASG.PNO} = \text{"P1"} \wedge \text{DUR} = 24)$$

# Query Decomposition

Analysis: enables rejection of normalized queries for which further processing is either impossible or unnecessary.

- A query is *type incorrect*: if any of its attribute or relation names are not defined in the global schema, or if operations are being applied to attributes of the wrong type.
- A query is *semantically incorrect* if its components do not contribute in any way to the generation of the result.

# Query Decomposition

Example: This query is type incorrect

```
SELECT    E#  
FROM      EMP  
WHERE     ENAME > 200;
```

For 2 reasons:

- Attribute E# is not declared in the schema
- the operation “>200” is incompatible with the type string of ENAME.

# Query Decomposition

- In the context of relational calculus, it is not possible to determine the semantic correctness of general queries.
- It is possible to do so based on the representation of the query as a graph, called a *query graph* or *connection graph*

# Query Decomposition

## *Query graph* or *connection graph*


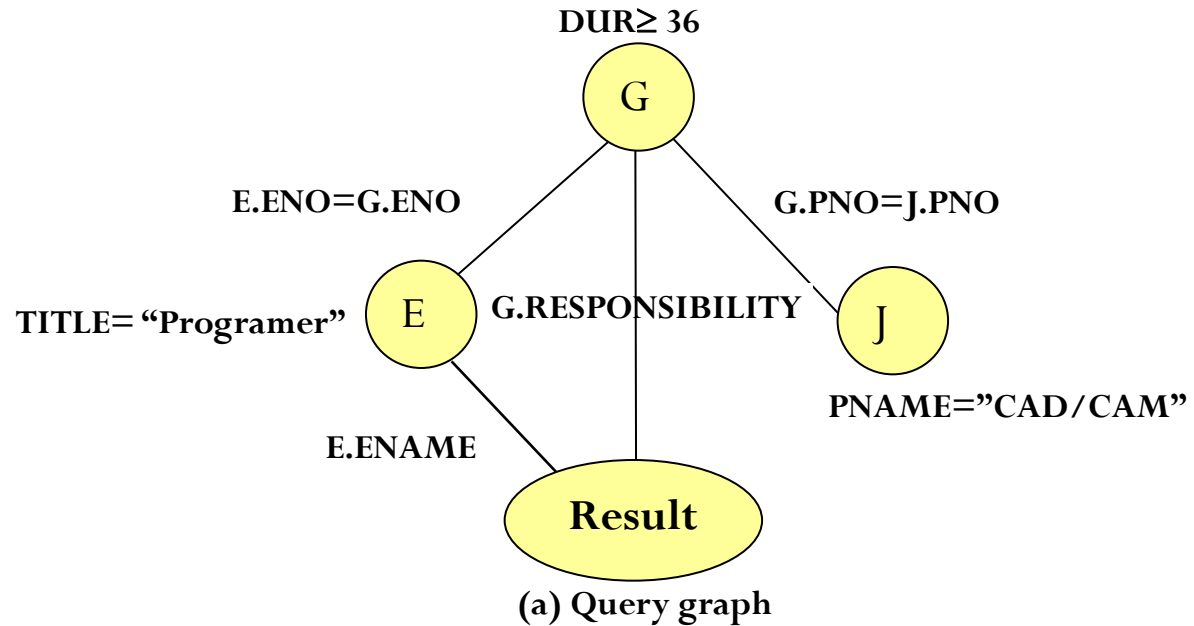
- One node indicates the result relation, and any other node indicates an operand relation
- An edge between two nodes one of which does not correspond to the result represents a join, whereas an edge whose destination node is the result represents a project
- A non-result node may be labeled by a select or a self-join (join of the relation with itself) predicate
- An important subgraph of the query graph is the *join graph*, in which only the joins are considered

# Query Decomposition

Example: “Find the names and responsibilities of programmers who have been working on the CAD/CAM project for more than 3 years.”

```
SELECT      ENAME, RESP
FROM        E, G, J
WHERE       E.ENO = G.ENO
AND         G.PNO = J.PNO
AND         PNAME = "CAD/CAM"
AND         DUR > 36
AND         TITLE = "Programmer";
```

# Query Decomposition





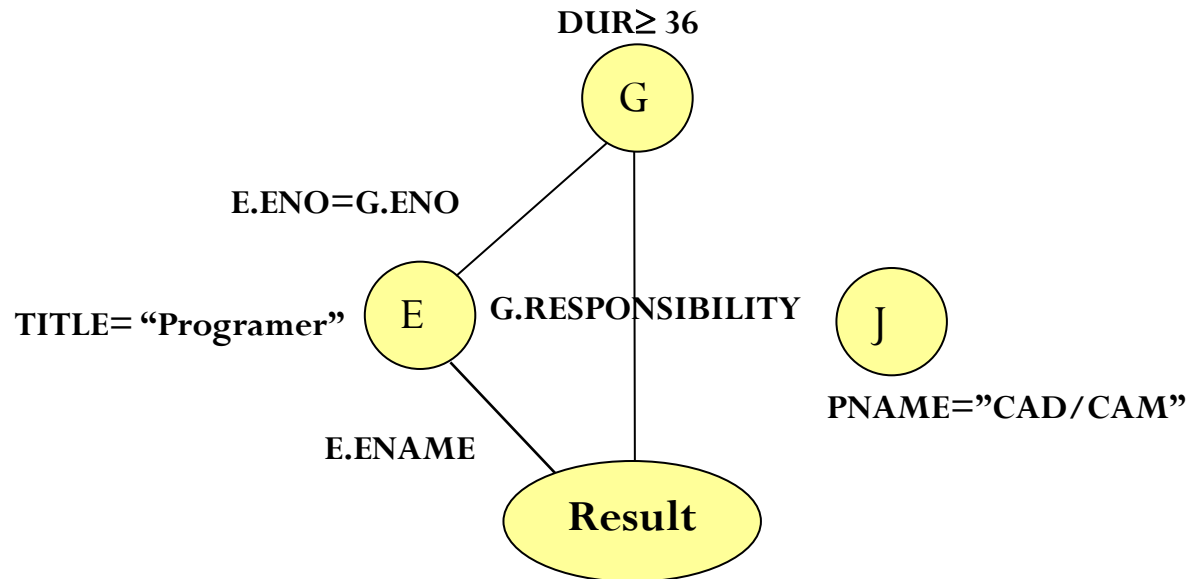
# Query Decomposition

Example: Consider the following query

```
SELECT      ENAME, RESP
FROM        E, G, J
WHERE       E.ENO = G.ENO
AND         PNAME = "CAD/CAM"
AND         DUR > 36
AND         TITLE = "Programmer";
```

# Query Decomposition

- Its query graph is disconnected, which tells us that the query is semantically incorrect.



Query graph

# Query Decomposition

Solutions to this problem:

- Reject the query
- Assume that there is an implicit Cartesian product between relations  $G$  and  $J$ ,
- Infer (using the schema) the missing join predicate  $G.PNO = J.PNO$  which transforms the query into that of previous Example.

# Query Decomposition

- Elimination of Redundancy: The enriched query qualification may contain *redundant predicates*. A naive evaluation of a qualification with redundancy can well lead to duplicated work. Such redundancy and thus redundant work may be eliminated by simplifying the qualification with the following well-known idempotency rules:

# Query Decomposition

$$1. p \wedge p \Leftrightarrow p$$

$$3. p \vee p \Leftrightarrow p$$

$$5. p \wedge \text{true} \Leftrightarrow p$$

$$7. p \vee \text{false} \Leftrightarrow p$$

$$9. p \wedge \text{false} \Leftrightarrow \text{false}$$

$$2. p \vee \text{true} \Leftrightarrow \text{true}$$

$$4. p \wedge \neg p \Leftrightarrow \text{false}$$

$$6. p \vee \neg p \Leftrightarrow \text{true}$$

$$8. p_1 \wedge (p_1 \vee p_2) \Leftrightarrow p_1$$

$$10. p_1 \vee (p_1 \wedge p_2) \Leftrightarrow p_1$$

# Query Decomposition

Example: the SQL query

```
SELECT      TITLE
FROM        E
WHERE (NOT (TITLE = "Programmer")
AND (TITLE = "Programmer" OR TITLE = "Elect. Eng.")
AND NOT (TITLE = "Elect. Eng.))
OR          ENAME = "J. Doe";
```

# Query Decomposition

Let:

$p_1$ : TITLE = “Programmer”,

$p_2$ : TITLE = “Elect. Eng.”,

$p_3$ : ENAME = “J. Doe”.

The query application is

$$(\neg p_1 \wedge (p_1 \vee p_2) \wedge \neg p_2) \vee p_3$$

$$\Leftrightarrow ((\neg p_1 \wedge p_1 \wedge \neg p_2) \vee (\neg p_1 \wedge p_2 \wedge \neg p_2)) \vee p_3$$

$$\Leftrightarrow ((\text{false} \wedge \neg p_2) \vee (\neg p_1 \wedge \text{false})) \vee p_3$$

$$\Leftrightarrow (\text{false} \vee \text{false}) \vee p_3$$

$$\Leftrightarrow p_3$$

# Query Decomposition

The query can be simplified using the previous rules to become

```
SELECT    TITLE
FROM      E
WHERE     ENAME = "J. Doe";
```



# Query Decomposition

Rewriting: rewrites the query in relational algebra. it is customary to represent the relational algebra query graphically by an *operator tree*.

- An operator tree is a tree in which a leaf node is a relation stored in the database, and a non-leaf node is an intermediate relation produced by a relational algebra operator.
- The sequence of operations is directed from the leaves to the root, which represents the answer to the query.

# Query Decomposition

The transformation of a tuple relational calculus query into an operator tree can easily be achieved as follows

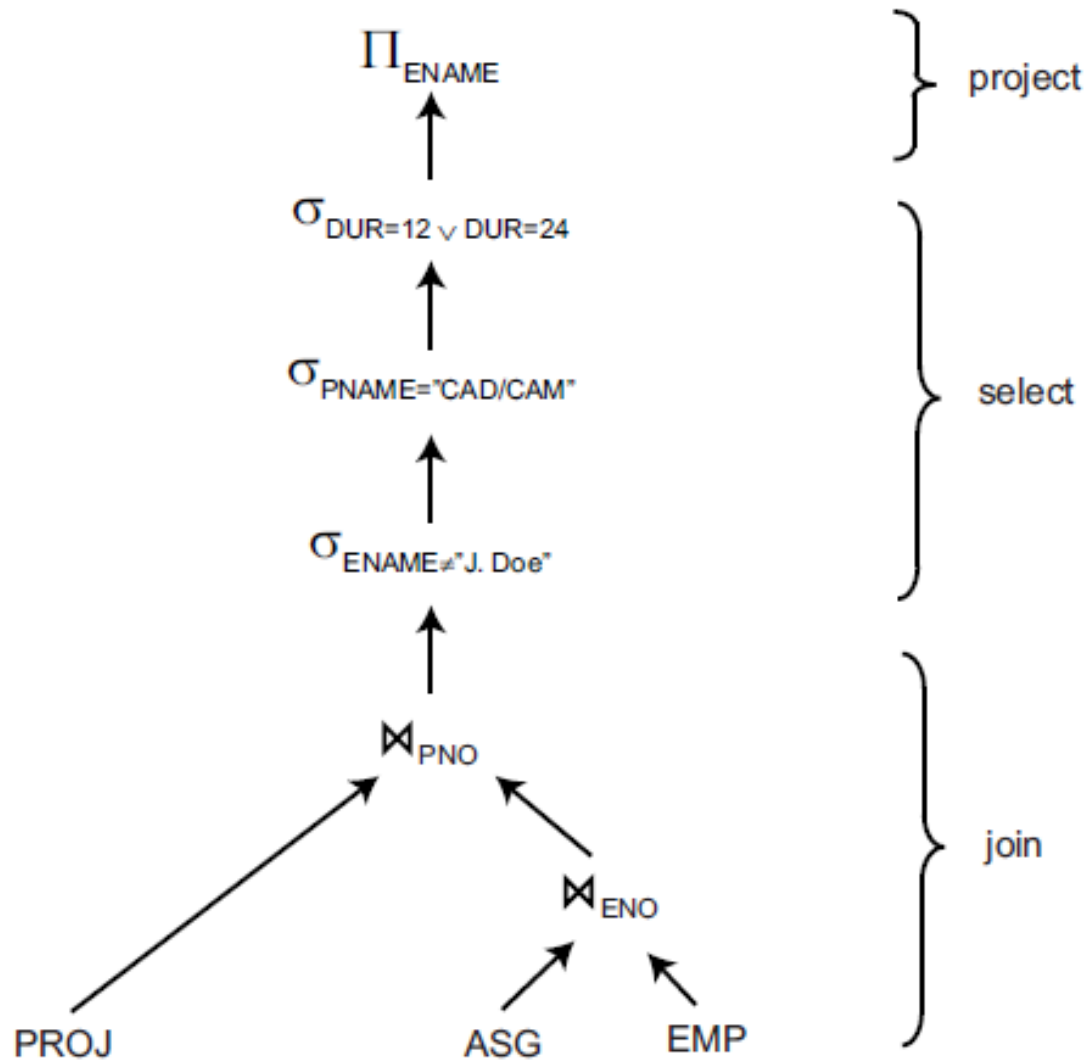
- A different leaf is created for each different tuple variable (corresponding to a relation). In SQL, the leaves are immediately available in the FROM clause.
- The root node is created as a project operation involving the result attributes. These are found in the SELECT clause in SQL.
- The qualification (SQL WHERE clause) is translated into the appropriate sequence of relational operations (select, join, union, etc.) going from the leaves to the root.

# Query Decomposition

Example: “Find the names of employees other than J. Doe who worked on the CAD/CAM project for either one or two years”

```
SELECT      ENAME
FROM        PROJ, ASG, EMP
WHERE       ASG.ENO = EMP.ENO
           AND   ASG.PNO = PROJ.PNO
           AND   ENAME != "J. Doe"
           AND   PROJ.PNAME = "CAD/CAM"
           AND   (DUR = 12 OR DUR = 24) ;
```

# Query Decomposition



# Query Decomposition

## 1. Commutativity of binary operators

$$R \times S \Leftrightarrow S \times R$$

$$R \bowtie S \Leftrightarrow S \bowtie R$$

## 2. Associativity of binary operators

$$R \times (S \times T) \Leftrightarrow (R \times S) \times T$$

$$R \bowtie (S \bowtie T) \Leftrightarrow (R \bowtie S) \bowtie T$$

## 3. Idempotence of unary operators

$$\Pi_{A'}(\Pi_{A''}(R)) \Leftrightarrow \Pi_{A'}(R), \text{ where } A', A'' \subseteq R \text{ and } A' \subseteq A''$$

$$\sigma_{p_1(A_1)}(\sigma_{p_2(A_2)}(R)) = \sigma_{p_1(A_1) \wedge p_2(A_2)}(R)$$

# Query Decomposition

## 4. Commuting selection with projection

$$\Pi_{A_1, \dots, A_n} (\sigma_{p(A_p)} (R)) = \Pi_{A_1, \dots, A_n} (\sigma_{p(A_p)} (\Pi_{A_1, \dots, A_n, A_p} (R)))$$

- Note that if  $A_p$  is already a member of  $\{A_1, \dots, A_n\}$ , the last projection on  $[A_1, \dots, A_n]$  on the right-hand side of the equality is useless.

$$\Pi_{A_1, \dots, A_n} (\sigma_{p(A_p)} (R)) = \sigma_{p(A_p)} (\Pi_{A_1, \dots, A_n, A_p} (R))$$

# Query Decomposition

## 5. Commuting selection with binary operators

$$\sigma_{p(A_p)}(R \times S) \Leftrightarrow \sigma_{p(A_p)}(R) \times S$$

$$\sigma_{p(A_i)}(R \bowtie_{p(A_i, B_k)} S) \Leftrightarrow \sigma_{p(A_i)}(R) \bowtie_{p(A_i, B_k)} S$$

$$\sigma_{p(A_i)}(R \cup T) \Leftrightarrow \sigma_{p(A_i)}(R) \cup \sigma_{p(A_i)}(T)$$

# Query Decomposition

## 6. Commuting projection with binary operators

- If  $C = A' \cup B'$ , where  $A' \subseteq A$ ,  $B' \subseteq B$ , and  $A, B$  are the sets of attributes over which relations  $R$  and  $S$ , respectively, are defined, we have

$$\Pi_C(R \times S) = \Pi_{A'}(R) \times \Pi_{B'}(S)$$

$$\Pi_C(R \bowtie_{p(A_i, B_j)} S) = \Pi_{A'}(R) \bowtie_{p(A_i, B_j)} \Pi_{B'}(S)$$

$$\Pi_C(R \cup S) = \Pi_{A'}(R) \cup \Pi_{B'}(S)$$




# Query Decomposition

These rules can be used in four different ways:

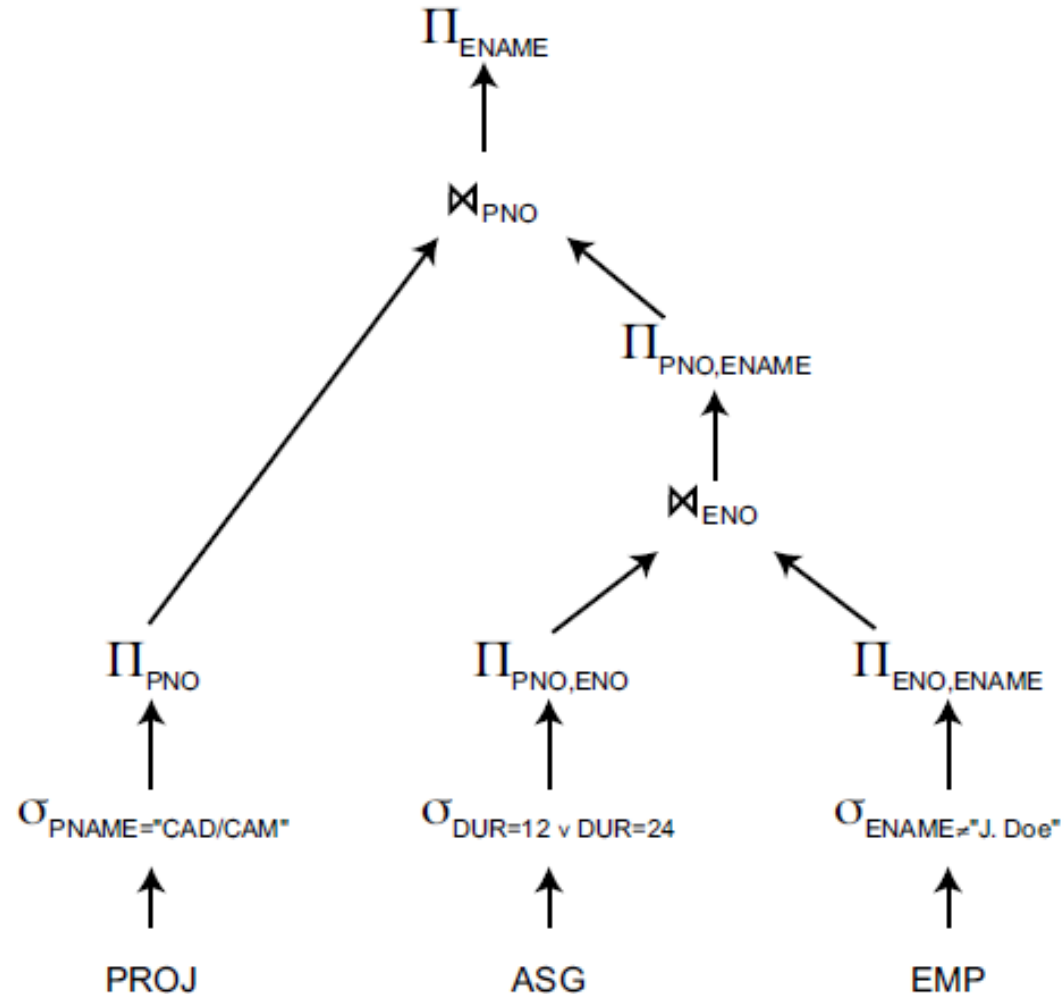
- First, they allow the separation of the unary operations, simplifying the query expression.
- Second, unary operations on the same relation may be grouped so that access to a relation for performing unary operations can be done only once.
- Third, unary operations can be commuted with binary operations so that some operations (e.g., selection) may be done first.
- Fourth, the binary operations can be ordered.

# Query Decomposition



Equivalent Operator Tree

# Query Decomposition



Rewritten Operator Tree

# Localization of Distributed Data

- General techniques for decomposing and restructuring queries apply to both centralized and distributed DBMSs and do not take into account the distribution of data.
- Localization layer translates an algebraic query on global relations into an algebraic query expressed on physical fragments
- Localization uses information stored in the fragment schema
- A global relation can be reconstructed by applying the reconstruction (or reverse fragmentation) rules and deriving a relational algebra program whose operands are the fragments. We call this a *localization program*.

# Localization of Distributed Data

- A naive way to localize a distributed query is to generate a query where each global relation is substituted by its localization program.
- **Localized query**: replacing the leaves of the operator tree of the distributed query with subtrees corresponding to the localization programs
- **Reduction techniques**: generate simpler and optimized queries by using the transformation rules and the heuristics, such as pushing unary operations down the tree

# Reduction for PHF

- The horizontal fragmentation function distributes a relation based on selection predicates

Example: Relation EMP(ENO, ENAME, TITLE) can be split into three horizontal fragments EMP1, EMP2, and EMP3

1.  $EMP1 = \sigma_{ENO \leq "E3"}(EMP)$
2.  $EMP2 = \sigma_{"E3" < ENO \leq "E6"}(EMP)$
3.  $EMP3 = \sigma_{ENO > "E6"}(EMP)$

- The localization program for an horizontally fragmented relation is the union of the fragments

$$EMP = EMP1 \cup EMP2 \cup EMP3$$

# Reduction with Selection

- Selections on fragments that have a qualification contradicting the qualification of the fragmentation rule generate empty relations.
- Given a relation  $R$  that has been horizontally fragmented as  $R_1, R_2, \dots, R_N$ , where  $R_i = \sigma_{p_i}(R)$ , the rule can be stated formally as follows:


- **Rule 1**:  $\sigma_{p_j}(R_i) = \emptyset$  if  $\forall x$  in  $R$ :  $\neg(p_i(x) \wedge p_j(x))$

where  $p_i$  and  $p_j$  are selection predicates,  $x$  denotes a tuple, and  $p(x)$  denotes “predicate  $p$  holds for  $x$ .”

# Reduction with Selection

- Example:

SELECT \*  
FROM EMP  
WHERE ENO = "E5";



(a) Localized query



(b) Reduced query



# Reduction with Join

- Joins on horizontally fragmented relations can be simplified when the joined relations are fragmented according to the join attribute.
- The simplification consists of distributing joins over unions and eliminating useless joins

$$(R_1 \cup R_2) \bowtie S = (R_1 \bowtie S) \cup (R_2 \bowtie S)$$

where  $R_i$  are fragments of  $R$  and  $S$  is a relation.

- **Rule 2**:  $R_i \bowtie R_j = \emptyset$  if  $\forall x$  in  $R_i, \forall y$  in  $R_j$  :  
 $\neg(p_i(x) \wedge p_j(y))$


fragments  $R_i$  and  $R_j$  are defined, respectively, according to predicates  $p_i$  and  $p_j$  on the same attribute

# Reduction with Join


- Example: Assume that relation EMP is fragmented between EMP1, EMP2, and EMP3, as above, and that relation ASG is fragmented as
  1.  $ASG1 = \sigma_{ENO \leq "E3"} (ASG)$
  2.  $ASG2 = \sigma_{ENO > "E3"} (ASG)$
- Consider the join query:

```
SELECT      *  
FROM        EMP, ASG  
WHERE       EMP.ENO = ASG.ENO;
```

# Reduction with Join



(a) Localized query



(b) Reduced query

# Reduction for Vertical Fragmentation

- The vertical fragmentation function distributes a relation based on projection attributes
- The localization program for a VF relation consists of the join of the fragments on the common attribute.
- Example: Relation EMP can be divided into two vertical fragments where the key attribute ENO is duplicated:

1.  $EMP1 = \Pi_{ENO,ENAME}(EMP)$

2.  $EMP2 = \Pi_{ENO,TITLE}(EMP)$

- The localization program is

$$EMP = EMP1 \bowtie_{ENO} EMP2$$

# Reduction for Vertical Fragmentation


- Queries on vertical fragments can be reduced by determining the useless intermediate relations and removing the subtrees that produce them
- Projections on a vertical fragment that has no attributes in common with the projection attributes (except the key of the relation) produce useless, though not empty relations.
- **Rule 3**:  $\Pi_{D,K}(R_i)$  is useless if the set of projection attributes  $D$  is not in  $A'$

Where relation  $R(A_1, A_2, \dots, A_n)$  is vertically fragmented as  $R_i = \Pi_{A'}(R)$ , and  $A' \subseteq \{A_1, A_2, \dots, A_n\}$


# Reduction for Vertical Fragmentation

- Example: consider the query

```
SELECT      ENAME
FROM        EMP
```



(a) Localized query



(b) Reduced query

# Reduction for Derived Fragmentation

- If relation  $R$  is subject to derived horizontal fragmentation due to relation  $S$ , the fragments of  $R$  and  $S$  that have the same join attribute values are located at the same site.
- $S$  can be fragmented according to a selection predicate.
- Derived fragmentation should be used only for one-to-many (hierarchical) relationships of the form  $S \rightarrow R$ , where a tuple of  $S$  can match with  $n$  tuples of  $R$ , but a tuple of  $R$  matches with exactly one tuple of  $S$ .

*Note that derived fragmentation could be used for many-to-many relationships provided that tuples of  $S$  (that match with  $n$  tuples of  $R$ ) are replicated.*

# Reduction for Derived Fragmentation

- Example: Given a one-to-many relationship from EMP to ASG, relation ASG(ENO, PNO, RESP, DUR) can be indirectly fragmented according to the following rules:

1.  $ASG1 = ASG \bowtie_{ENO} EMP1$

2.  $ASG2 = ASG \bowtie_{ENO} EMP2$

Where:  $EMP1 = \sigma_{TITLE="Programmer"}(EMP)$

$EMP2 = \sigma_{TITLE \neq "Programmer"}(EMP)$

- The localization program for a horizontally fragmented relation is the union of the fragments

$$ASG = ASG1 \cup ASG2$$



# Reduction for Derived Fragmentation

- Queries on derived fragments can also be reduced: certain joins will produce empty relations if the fragmentation predicates conflict.

- Example: the predicates of ASG1 and EMP2 conflict, thus

$$\text{ASG1} \bowtie \text{EMP2} = \emptyset$$


- The reduced query is always preferable to the localized query because the number of partial joins usually equals the number of fragments of R.

# Reduction for Derived Fragmentation

- Example:


```

SELECT      *
FROM        EMP, ASG
WHERE       ASG.ENO = EMP.ENO
AND        TITLE = "Mech. Eng."
    
```




(a) Localized query

# Reduction for Derived Fragmentation




(b) Query after pushing selection down



(c) Query after moving unions up

# Reduction for Derived Fragmentation



(d) Reduced query after eliminating the left subtree

# Reduction for Hybrid Fragmentation

- The goal of hybrid fragmentation is to efficiently support queries involving projection, selection, and join.
- The optimization of an operation or of a combination of operations is always done at the expense of other operations.
- The localization program for a hybrid fragmented relation uses unions and joins of fragments

# Reduction for Hybrid Fragmentation

- Example:

1.  $EMP1 = \sigma_{ENO \leq "E4"}(\Pi_{ENO, ENAME}(EMP))$

2.  $EMP2 = \sigma_{ENO > "E4"}(\Pi_{ENO, ENAME}(EMP))$

3.  $EMP3 = \Pi_{ENO, TITLE}(EMP)$

The localization program is

$$EMP = (EMP1 \cup EMP2) \bowtie_{ENO} EMP3$$



# Reduction for Hybrid Fragmentation

- Queries on hybrid fragments can be reduced by combining the rules:
  1. Remove empty relations generated by contradicting selections on horizontal fragments.
  2. Remove useless relations generated by projections on vertical fragments.
  3. Distribute joins over unions in order to isolate and remove useless joins.

# Reduction for Hybrid Fragmentation

- Example: application of rules (1) and (2)

SELECT        ENAME  
FROM         EMP  
WHERE        ENO="E5"





# Query Optimization in DDBS

- Query decomposition and data localization are the two successive functions that map a calculus query, expressed on distributed relations, into an algebraic query (query decomposition), expressed on relation fragments (data localization).
- Query decomposition can generate an algebraic query simply by translating into relational operations the predicates and the target statement as they appear.
- Data localization can, in turn, express this algebraic query on relation fragments, by substituting for each distributed relation an algebraic query corresponding to its fragmentation rules.